

BI Metric Governance Handbook

Document control	Value
Document type	Controlled analytics policy
Owner	Analytics Platform - Metric Governance
Approved by	BI Governance Council, Marketplace Operations, Data Stewardship
Audience	Finance, operations, marketing, seller success, BI analysts, and BI-agent users
Effective date	2026-06-01
Review cadence	Quarterly, or within 10 business days of schema or policy changes
Source basis	Public Olist Brazilian E-Commerce dataset and Marketing Funnel dataset descriptions, plus synthetic governance rules for a BI-agent evaluation corpus.
Confidentiality	Internal use only - sandbox corpus; do not represent as official Olist policy

RAG ingestion note: This document is intentionally authored as selectable text, with stable headings, tables, examples, policy precedence notes, and retrieval keywords. Chunk at section boundaries; preserve tables as Markdown or CSV in the vector index.

Version History

Version	Date	Author / owner	Summary of change
1.0	2026-06-01	AI sandbox generator	Initial two-page demo document.
2.0	2026-06-06	BI Knowledge Management	Expanded into searchable internal-style policy with controls, examples, decision rules, and RAG guidance.

Purpose

This handbook defines approved business metrics for the governed BI copilot. The goal is to make answers consistent when users ask about sales, orders, delivery, reviews, seller performance, and marketing funnel conversion. The copilot must retrieve this handbook before generating SQL for any governed metric.

Policy Precedence

Priority	Rule source	How the BI agent must use it
1	User role and data-access policy	Apply access controls before producing SQL, rows, exports, or traces.
2	Metric governance handbook	Use approved metric definitions and caveats before generating SQL.
3	Domain playbooks and SLA policies	Use thresholds and interpretation rules only after metrics have been calculated.
4	User preference	Honor display preferences only when they do not conflict with controls.

Dataset Scope and Table Contracts

The public e-commerce dataset includes order, item, payment, customer, seller, product, geolocation, and review information. The marketing funnel dataset includes marketing qualified leads and closed deals that can be linked into seller acquisition analysis. In this sandbox, table names use the conventional Olist CSV naming pattern.

Business area	Canonical table names	Primary join keys	Default date
Orders	olist_orders_dataset, olist_order_items_dataset	order_id, customer_id, seller_id	order_purchase_timestamp
Payments	olist_order_payments_dataset	order_id	payment_sequential when needed; otherwise order date

Business area	Canonical table names	Primary join keys	Default date
Reviews	olist_order_reviews_dataset	order_id, review_id	review_creation_date
Products and sellers	olist_products_dataset, olist_sellers_dataset	product_id, seller_id	Use order date from joined orders
Customers and geography	olist_customers_dataset, olist_geolocation_dataset	customer_id, customer_zip_code_prefix	Use order date from joined orders
Marketing funnel	olist_marketing_qualified_leads_dataset, olist_closed_deals_dataset	mql_id, seller_id	first_contact_date for MQL cohorts; won_date for closing analysis

Global Rules

- Revenue means product sales value from delivered orders, excluding freight.
- GMV means product sales plus freight from delivered orders.
- Paid amount means customer payment value and should use payment data, not item price.
- Order counts must use distinct order IDs after joins.
- Review metrics are order-level and must not be duplicated across item-level joins.
- Marketing funnel conversion must use marketing qualified leads as the base population.
- Date filters should use the default date column for the selected metric unless the user asks otherwise.
- Answers should include caveats when a metric excludes freight, canceled orders, unconverted leads, or missing delivery dates.

Approved Metrics

Metric	Business definition	Formula / required logic	Default filters and caveats
Revenue	Total product sales value from delivered marketplace orders.	SUM(order_items.price) after joining orders to order_items.	orders.order_status = delivered. Excludes freight, canceled, unavailable, shipped, invoiced, and non-delivered orders. Do not use payment_value unless user asks for paid amount.
Gross Merchandise Value	Product value plus freight value for delivered orders.	SUM(order_items.price + order_items.freight_value).	Includes freight. Does not account for refunds, discounts, commissions, or platform fees.
Order Count	Number of distinct marketplace orders.	COUNT(DISTINCT orders.order_id).	Always deduplicate order IDs after joining item, payment, review, product, or seller tables.
Average Order Value	Average product revenue per delivered order.	First aggregate item price to order level; then divide by distinct delivered orders.	Do not use AVG(order_items.price), because that produces average item price. Excludes freight unless GMV-based AOV is requested.
Late Delivery Rate	Percentage of delivered orders where customer delivery happened after estimated delivery date.	Late delivered orders divided by delivered orders with known customer delivery date.	Use customer delivery date, not carrier handoff date. Do not include undelivered orders in the denominator.
Average Review Score	Average customer review score for reviewed orders.	AVG(order_reviews.review_score) at order level.	Reviews are order-level. Pre-aggregate at order level before category or seller analysis.
Lead Conversion Rate	Percentage of MQLs that became closed deals.	COUNT(DISTINCT closed_deals.mql_id) / COUNT(DISTINCT marketing_qualified_leads.mql_id).	Use marketing qualified leads as the base table and left join closed deals. Do not use closed deals as denominator.

SQL Reference Patterns

Question type	SQL pattern that should be generated	Validation checks
Revenue by month	Join orders to order_items; filter delivered; group by DATE_TRUNC(month, order_purchase_timestamp).	No freight in revenue; no non-delivered orders; no payment_value.
Late delivery by state	Use orders joined to customers; require delivered and non-null customer delivery date; group by customer_state.	Report delivered order count; avoid ranking states with small volume unless requested.
Seller review performance	Join reviews to orders and items; pre-aggregate review score at order level before seller attribution.	Flag multi-seller orders as directional if attribution is ambiguous.
MQL conversion by origin	Start with marketing_qualified_leads; left join closed_deals; group by origin.	Include MQL count and closed deal count; warn when MQL count < 30.

Copilot Behavior

- Prefer approved metrics over ad hoc calculations.
- Ask clarification when “sales” could mean revenue, GMV, paid amount, or order count.
- Reject or warn about SQL that violates approved metric rules.
- Include metric name, default filters, and material caveats in final answers when useful.
- Every numerical claim must be tied to database results and, when relevant, retrieved policy context.

RAG Retrieval Tests

Test prompt	Expected retrieved sections	Expected answer behavior
What does revenue mean in this BI agent?	Metric definitions; caveats; default filters	Use approved definition, cite default filters, ask clarification if user means GMV or paid amount.
Can a viewer see SQL or row-level seller details?	Access levels; SQL visibility; sensitive data rules	Deny raw detail, provide aggregate summary, explain role limitation.
Which sellers or regions need attention this month?	SLA thresholds; seller performance rules; incident triggers	Compute metrics first, apply thresholds, include volume and caveats.

Document End

For production use, replace sandbox owners and thresholds with approved policy owners, legal review, actual access-control groups, and validated database schema contracts.